

Nicholas Clarke, PhD

Independent Researcher - AI Alignment, Moral Philosophy & Evaluation
Cambridge, UK | nicholasclarke.github.io | nclarke431@gmail.com

RESEARCH PROFILE

Independent philosopher and AI-alignment researcher working on alignment-target specification, value robustness under optimisation, and AI-assisted moral discovery. My central research programme investigates whether idealised human preferences - what people would endorse under improved information, rational deliberation, and greater coherence - can provide more defensible and empirically tractable alignment targets than uncorrected behavioural preferences. The work connects moral philosophy, behavioural science, AI safety, evaluation, and governance, with particular attention to proxy failure, Goodhart effects, deception, corrigibility, pluralism, and human oversight.

CURRENT RESEARCH PROGRAMME

- **AI-Assisted Discovery of Idealised Human Desire: Toward a Naturalistic Science of Moral Alignment** - research programme reframing alignment as a problem of target discovery as well as behavioural control. Develops an empirical pathway for studying value transformation under improved epistemic and deliberative conditions, with a corrigibility-first architecture that keeps deployment authority with humans.
- **Planned empirical work:** operational measures of information sufficiency, deliberative quality, and coherence; structured idealisation interventions; longitudinal preference-stability analysis; modelling of convergence and persistent disagreement; and red-team evaluation of manipulation and proxy failure.
- **Funding and collaboration:** actively seeking support and research partnerships for an exploratory pilot phase and broader interdisciplinary development.

SELECTED RESEARCH & WRITING

- **Beyond Behavioural Control: Idealised Human Desire as an Alignment Target for Robust AI Alignment** - Working paper. Develops the distinction between behavioural/control mechanisms and the prior alignment-target problem; proposes idealised human desire as a revisable candidate target and reframes alignment as moral discovery.
- **Humanist AI Needs a Theory of What Alignment Is For** - Public response to Microsoft AI's 2026 Humanist AI Code of Conduct consultation. Argues for explicit alignment-target discovery, empirically informed idealisation, pluralism-sensitive evaluation, and epistemic integrity as a safety property.
- **Goal Theory** - Ongoing programme in naturalistic moral philosophy grounded in idealised human desire, with work on moral disagreement, moral luck, non-identity, and metaethical semantics.
- **Epistemology and reasons under non-ideal conditions** - Working papers developing non-factive operational accounts of reasons and knowledge suited to deliberation under uncertainty and fallible cognition.

EDUCATION

- **PhD in Philosophy - University of London** | Awarded with no corrections. Doctoral research developed and defended Goal Theory, a naturalistic account of morality centred on idealised human desire.
- **MA in Philosophy (Distinction) - University of London** | Dissertation applying Bayesian reasoning to complex evidential claims under uncertainty.

ACADEMIC & PUBLIC ENGAGEMENT

- **Associate Lecturer - AI Ethics & Technology, Anglia Ruskin University** | Research-informed teaching on AI ethics, risk, alignment, and governance.
- **Research seminar on AI alignment, Anglia Ruskin University (2026)** | Presented the AI-assisted moral-discovery research programme and discussed philosophical and practical challenges of alignment-target specification.
- **Public writing** | Essays on AI alignment, ethics, moral philosophy, and rational inquiry on Substack and personal research site.

PROFESSIONAL EXPERIENCE

- **Technical Author - Darktrace** | Technical documentation and structured content work for complex cybersecurity products, including documentation review, migration, clarity/consistency improvement, and collaboration with technical stakeholders.
- **Reporting & Information Analyst - University of Cambridge** | Empirical analysis, structured evaluation, and decision-focused reporting across large organisational datasets.
- **Earlier career** | Software development, technical writing, analytics, and senior roles in enterprise systems and SaaS environments.

METHODS & SKILLS

Conceptual analysis • research design • evidence synthesis • experimental design • human-in-the-loop studies • quantitative analysis • value/preference modelling • failure-mode analysis • systems thinking • technical writing • research communication • interdisciplinary collaboration

RESEARCH THEMES

AI alignment and safety • alignment-target specification • idealised human preferences • empirical moral discovery • corrigibility and human control • deception and proxy failure • Goodhart effects • moral uncertainty and pluralism • naturalistic metaethics • epistemology under non-ideal conditions

ONLINE

- Research profile and selected writing: nicholasclarke.github.io
- Microsoft consultation response: [Humanist AI Needs a Theory of What Alignment Is For](#)